

# Unsupervised and Unstructured Machine Learning

BA820 – Mohannad Elhamod

# Midterm Review

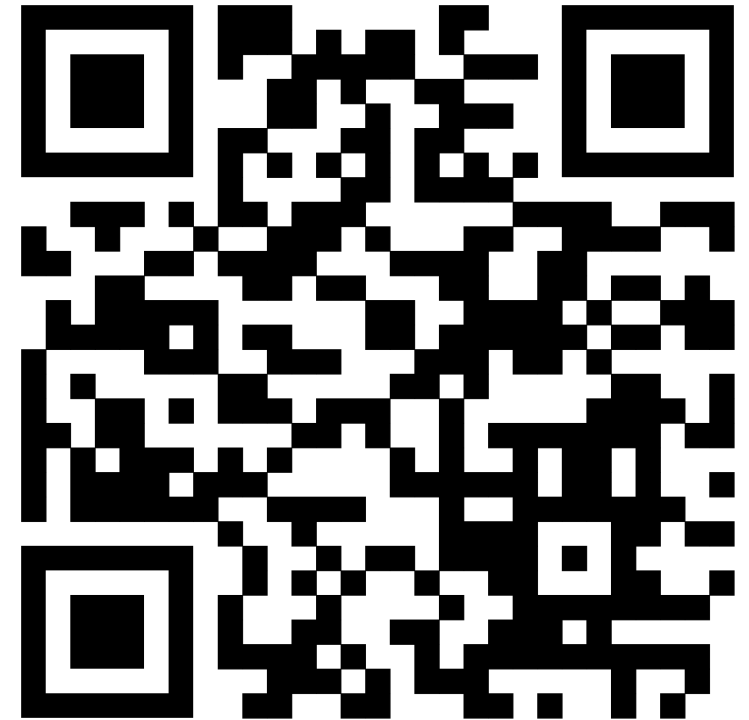


# Start Stop Continue

---

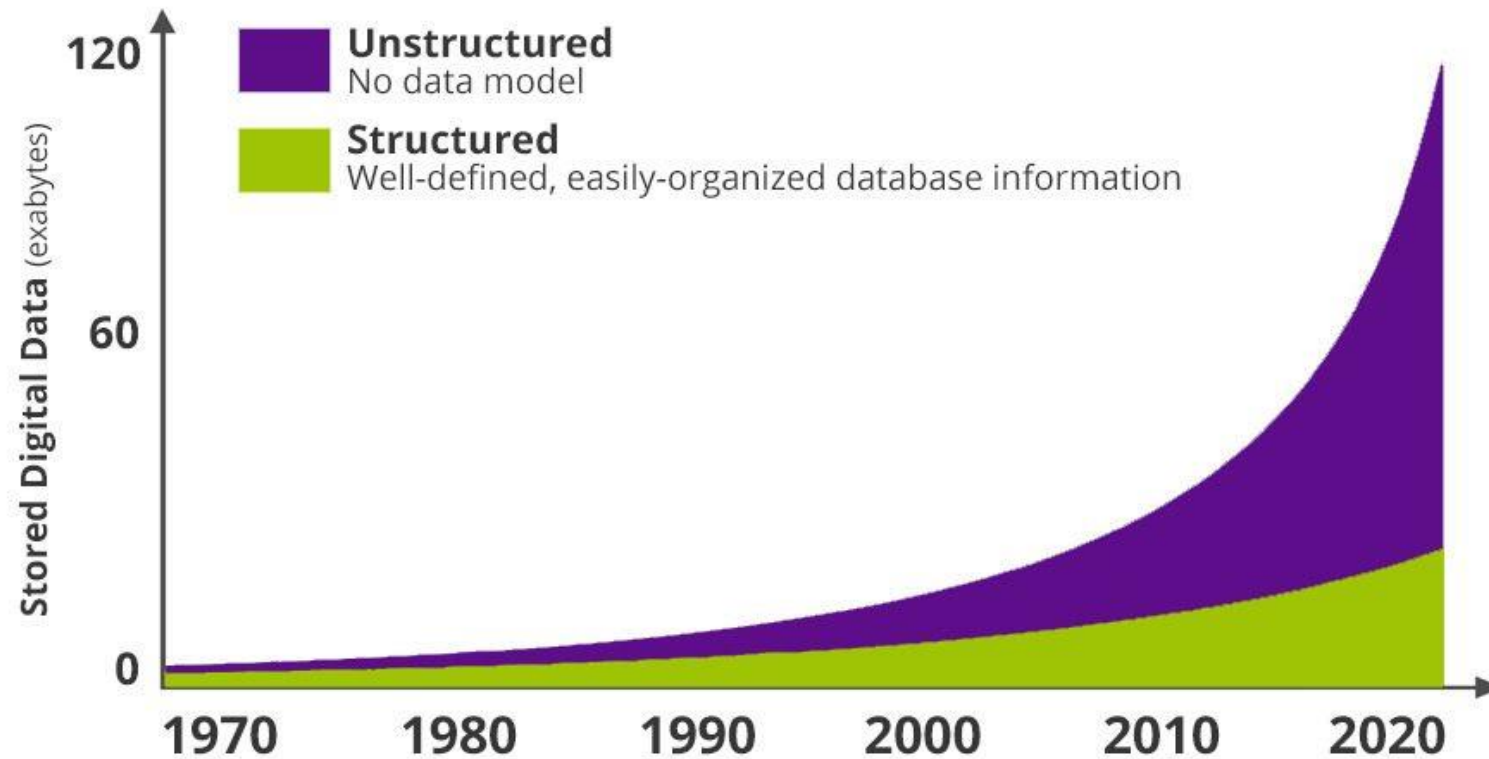
# Please fill this out

- <https://forms.gle/VuHJUjtzLrrr1TNY8>



# Text Mining

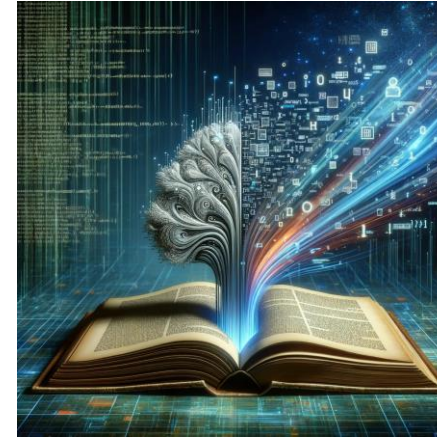
# Most Data is No Longer Structured...



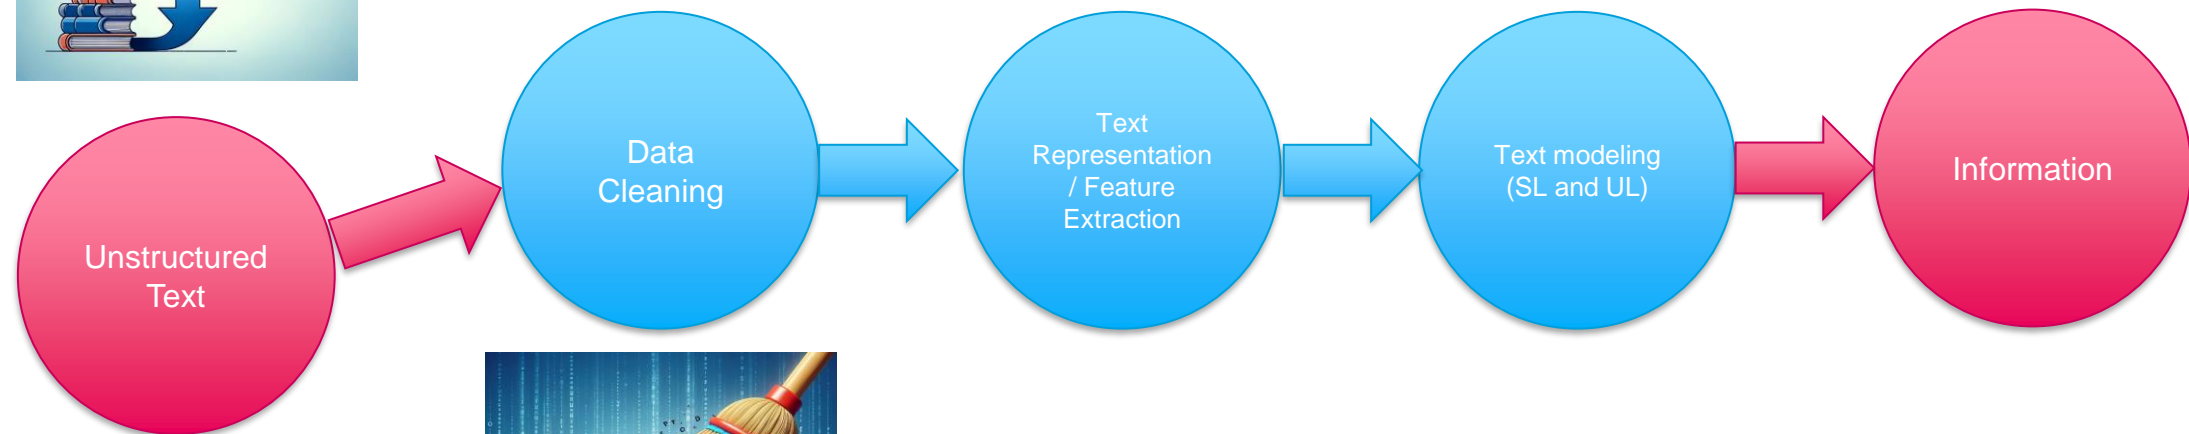


<https://www.komprise.com/>

# What can we do with this data?

- Text classification
- Text generation
- Text summarization
- Music recommendation
- Image categorization
- ....



# Text Mining





# Collecting Unstructured Data

- How can we extract this data?
  - Datasets found publicly (UCI, Kaggle, [Common Crawl](#), [Wikimedia Downloads](#), etc.).
  - Using APIs (e.g., [twitter API](#), [News API](#)).
  - Web scrapping (e.g., [BeautifulSoup](#), [Selenium](#))
  - Own private data.



# Data Cleaning with *Regex*

- Just like structured/tabular data, we generally need to clean up the text to make it more useful.
  - Examples: case-sensitivity, punctuation, etc.
- *Regex* is used for finding string matches and formatting text:
  - Playground:  
[https://www.w3schools.com/python/python\\_regex.asp](https://www.w3schools.com/python/python_regex.asp)
  - Cheat sheet:  
<https://www.debuggex.com/cheatsheet/regex/python>

# How to Represent Text?

- Computers do not understand text...
- We need to represent text in a language they understand... numbers!
- Simple proposals (we are just brainstorming here):
  - Each sentence is represented in terms of the words it contains...
    - This is called Tokenization.
  - Each word is represented by a number...
    - This is called Vectorization.

# Tokenization: Some Terminology

- **Document:** A body of text (e.g., a tweet, a pdf, an article, etc.).
- **A token:** The building block of a document.
  - Examples: character-level, word-level, ...
- **A separator:** Special tokens that split a document into tokens.
  - Examples: punctuation, spaces,
- Demo

# Need for Advanced Text Pre-Processing

- Some tokens may occur too frequently in any text without contributing much to its meaning. These are called stop words and are generally removed.
- Other issues:
  - **Stemming**: big, bigger, biggest
  - **Lemmatization**: drive drove driven
  - **Homonyms**: bank (river or money?)
  - **Synonym**: Yes, sure.
- We will come back to all of this later...

# Syntax vs. Semantics

- Notice that:
  - Different tokens might have the same meaning (e.g., like, enjoy)
  - The same token might have different meanings (is *like* something, to *like* something)
- **So, while tokens represent syntax, we really care about meaning/semantics.**
- In many cases, you can only get the meaning through **context** (i.e., the token's place with respect to other tokens.)
- We will come back to all of this later...

# Text Modeling

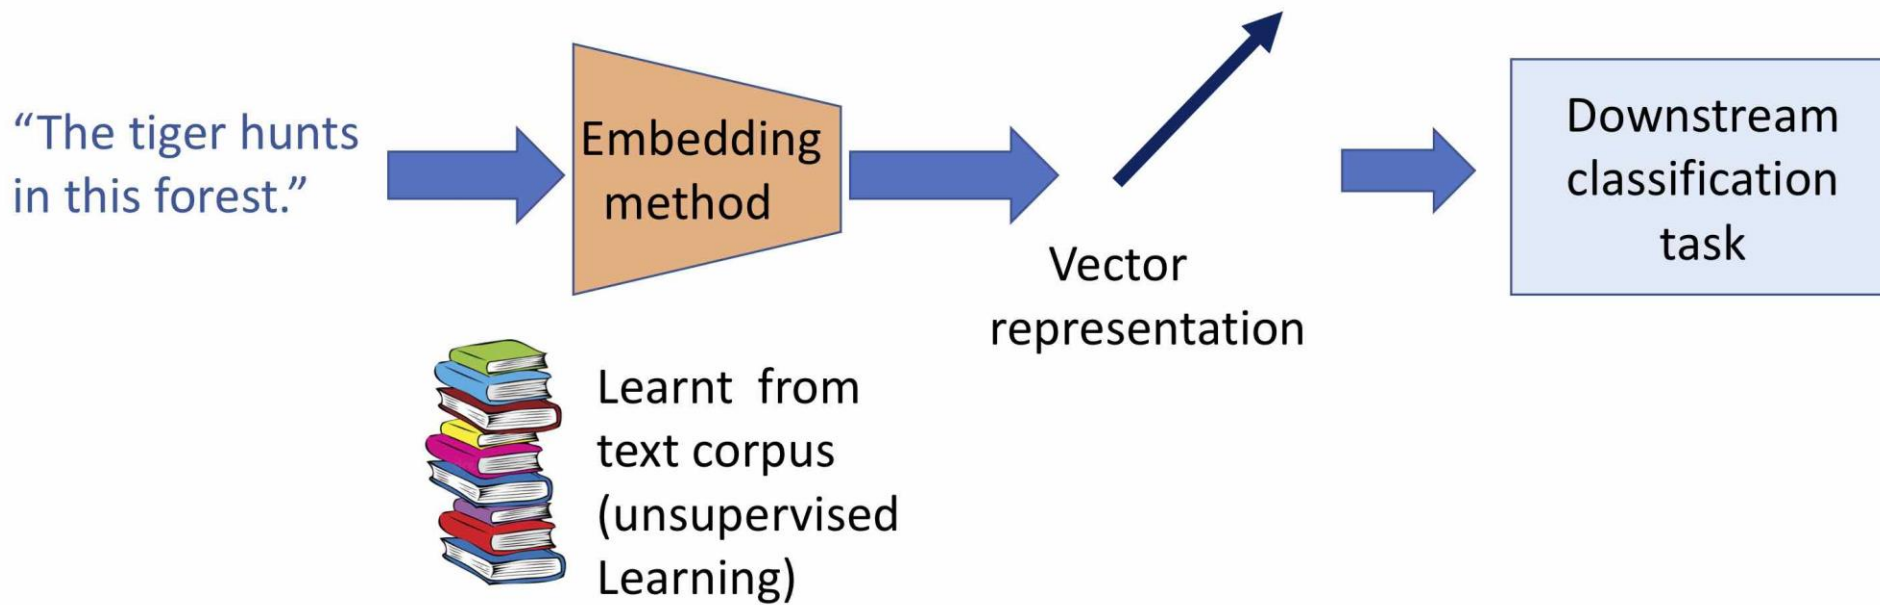
- Once text data is in the proper representation (i.e., tokenized and vectorized), we can apply the methods we have learned so far:
  - Unsupervised ML (e.g., dimensionality reduction, clustering, etc.).
  - Supervised ML (e.g., classification, translation, etc.).

# Text Vectorization



# Vectorization: Text as Numbers

Typical pipeline for unsupervised text embedding



[Offconvex.org](https://offconvex.org)

# How Would You Represent a Document?

- Let's start simple. We could represent a document simply as a collection of tokens.
  - Vectorization by document (not token).
- This approach is called Bag of Words (BoW).
- Demo**

Sentence	hockey	fun	i	like	golf
I like golf!			1	1	1
I like hockey.	1		1	1	
Hockey and golf are fun!	1	1			1

# Bag of Words (BoW)

- **Cons:**

- Disregards word order!
- Number of features can be exhaustive.
- **Frequency bias.**
  - (e.g., If the word “space” appears in a children’s book, it carries more significance than when it appears in an article about galaxies.

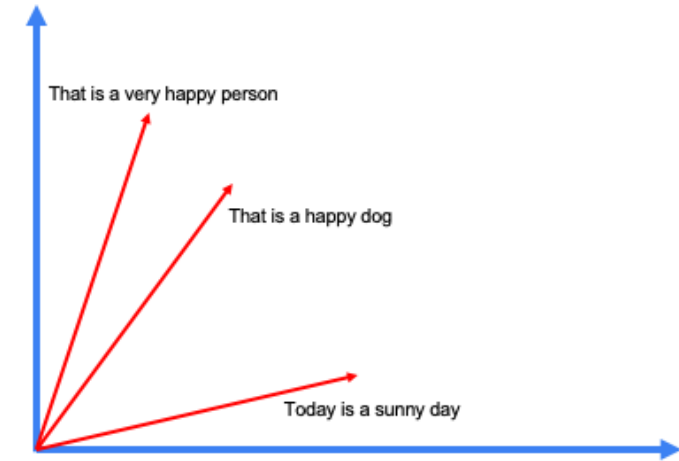
- **Pros:**

- Simple to implement

	I	hate	love	golf	soccer
I hate golf and love soccer	1	1	1	1	1
I hate soccer and love golf	1	1	1	1	1

# Document Similarity

- How do we measure if two documents are similar?
  - We need a metric like Euclidean distance.
  - But... What if documents have different lengths?
- We need a metric that is robust to differences in document size...
  - Enter Cosine Similarity.

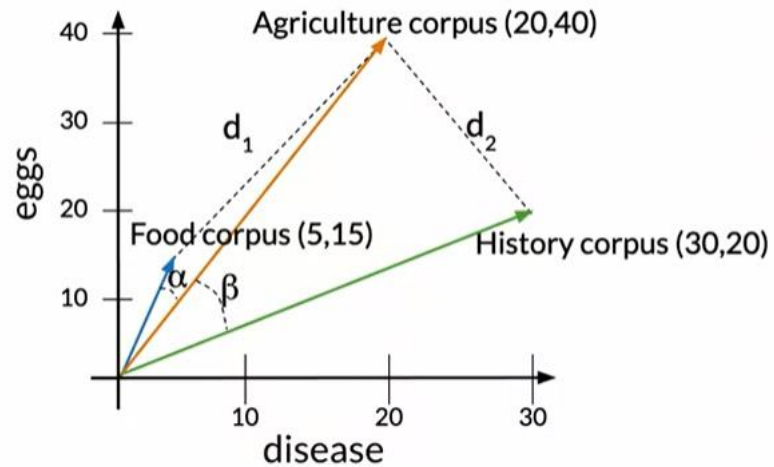


<https://mlops.community>

$$\text{similarity} = \cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|} = \frac{\sum_{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n A_i^2} \sqrt{\sum_{i=1}^n B_i^2}}$$

# Document Similarity

## Euclidean distance vs Cosine similarity



The cosine of the angle  
between the vectors

[medium.com](https://medium.com/@davidfriesen/euclidean-distance-vs-cosine-similarity-400000000000)

$$\text{similarity} = \cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|} = \frac{\sum_{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n A_i^2} \sqrt{\sum_{i=1}^n B_i^2}}$$

# Solving for Frequency Bias: TF-IDF

- Term Frequency – Inverse Document Frequency
- Instead of simply counting the absolute number of occurrences, how about we adjust in terms of token presence across documents?

## TF-IDF

Term  $x$  within document  $y$

$tf_{x,y}$  = frequency of  $x$  in  $y$

$df_x$  = number of documents containing  $x$

$N$  = total number of documents

$$w_{x,y} = tf_{x,y}$$

BoW frequency

→ Portion of documents the token appears in

# Solving for Frequency Bias: TF-IDF

- The more commonly a word exists across documents the less important it is!
- What happens when a term (i.e., token) appears in all documents?
- **Demo**

$$w_{x,y} = \text{tf}_{x,y} \times \log \left( \frac{N}{df_x} \right)$$

# Context Matters...

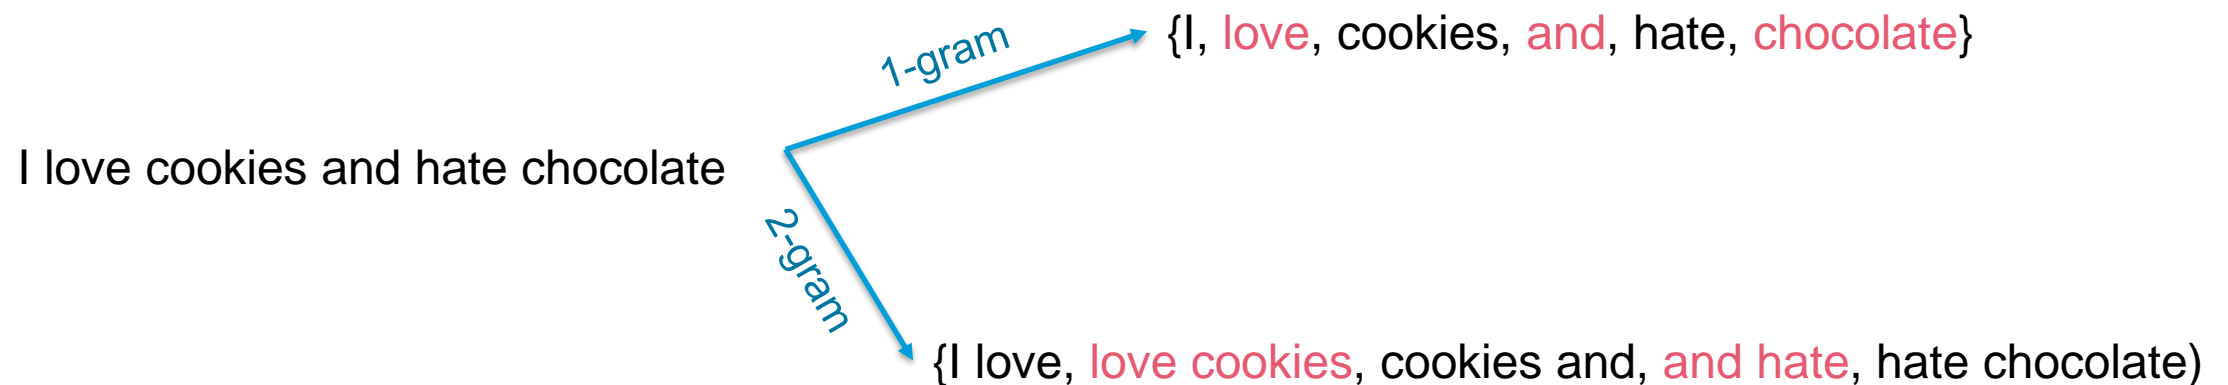


# So far...

- We calculated a score for each individual token (e.g., the frequency of the word in BoW).
- Have our methods captured context so far?
  - Context has so far been captured as *“the collection of words that appear together in a document”*, regardless of order.

# Can We Do Better?

- Tokens that are immediately next to each other are generally more related to each other.
- Can we apply this concept at the tokenization level?



# ***n***-grams

- An  $n$ -gram is a sequence of  $n$  items (i.e., tokens).
- $n$  refers to the context size.
- The larger  $n$  is, the more context a token captures.
- What happens to the number of possible tokens as  $n$  increases?
- What happens to the vector representation as  $n$  increases?

# How Much Context Is Needed?

- A large  $n$  captures more context, but it's more restrictive in terms of finding matches.
- A small  $n$  captures less context. but a token is more likely to be matched in other documents.
- Maybe we can include both? Any draw backs?
- **Demo**



# Meaning: The Illusive Goal...

# Where Does Meaning Come From?

I went to the **branch** and deposited some money.

I went to the **bank** and deposited some money.

I went to the **ATM** and deposited some money.

Words which frequently appear in **similar contexts** have **similar meaning**.

[Lena-votta](#)

# Representing Meaning as a Vector

- So, why don't we construct vector representations such that words that appear in the same context have similar vectors.
  - embeddings = vector representation = features = latent space.
  - $v_1$  and  $v_2$  are similar  $\Rightarrow \|v_1 - v_2\| \approx 0$
- [Example from images.](#)

# Word Embeddings

- Represent each word as a vector of some pre-defined size/dimensionality.
- You can visualize these embeddings using PCA or t-SNE.
- Some interesting properties will appear:
  - [Demo 1 \(directionality\)](#)
  - [Demo 2 \(semantics\)](#)




# How are such embeddings constructed?

- For each token, start with a randomly initialized vector.
- From your training text, construct a table showing which tokens co-occur and which don't.
- For words that co-occur, push the vectors towards each other.
- For words that don't, push them away from each other.
- Examples of word embeddings:
  - *Word2Vec*
  - *GloVe*

# Something is still missing...

- Well, now we have word embeddings...
  - But, how do we get sentence embeddings?!
- We need to somehow aggregate the word embeddings.
  - There are many ways to do this...
  - The most basic one is taking the average embedding.
  - Another way is to take the TF-IDF weighted averaging.
  - Or.... Using a Neural Net. Welcome to Deep Learning!



# So, what's the big deal?

Demo (Predict the next word)

# Demo

# Language Modeling

- A playground for predicting the next word:

<https://pair.withgoogle.com/explorables/fill-in-the-blank/>